

Three function of Virtual Memory

◦ Translation

- A program is given consistent view of memory, thus a programmer can write a program without knowing the physical layout of memory.
- Address space can be continuous while it can grow
- Uses table lookup to translate virtual address into physical address
- Each process has its own address space
=> Each process has its own mapping.

◦ Protection

- Different processes protected from each other by having different mapping between different processes.
- Kernel data protected from User programs

◦ Sharing

- Two processes can have a page table entry that is mapped to the same physical page.



Virtual Memory

Translation:

- Program can be given consistent view of memory (virtual memory), no matter how the physical memory is organized.
- Only the most important part of program (“Working Set”) must be in physical memory
- Contiguous structures (like stacks) use only as much physical memory as necessary yet still grow later
- Uses table lookup to translate virtual address into physical address



Virtual Memory

Protection:

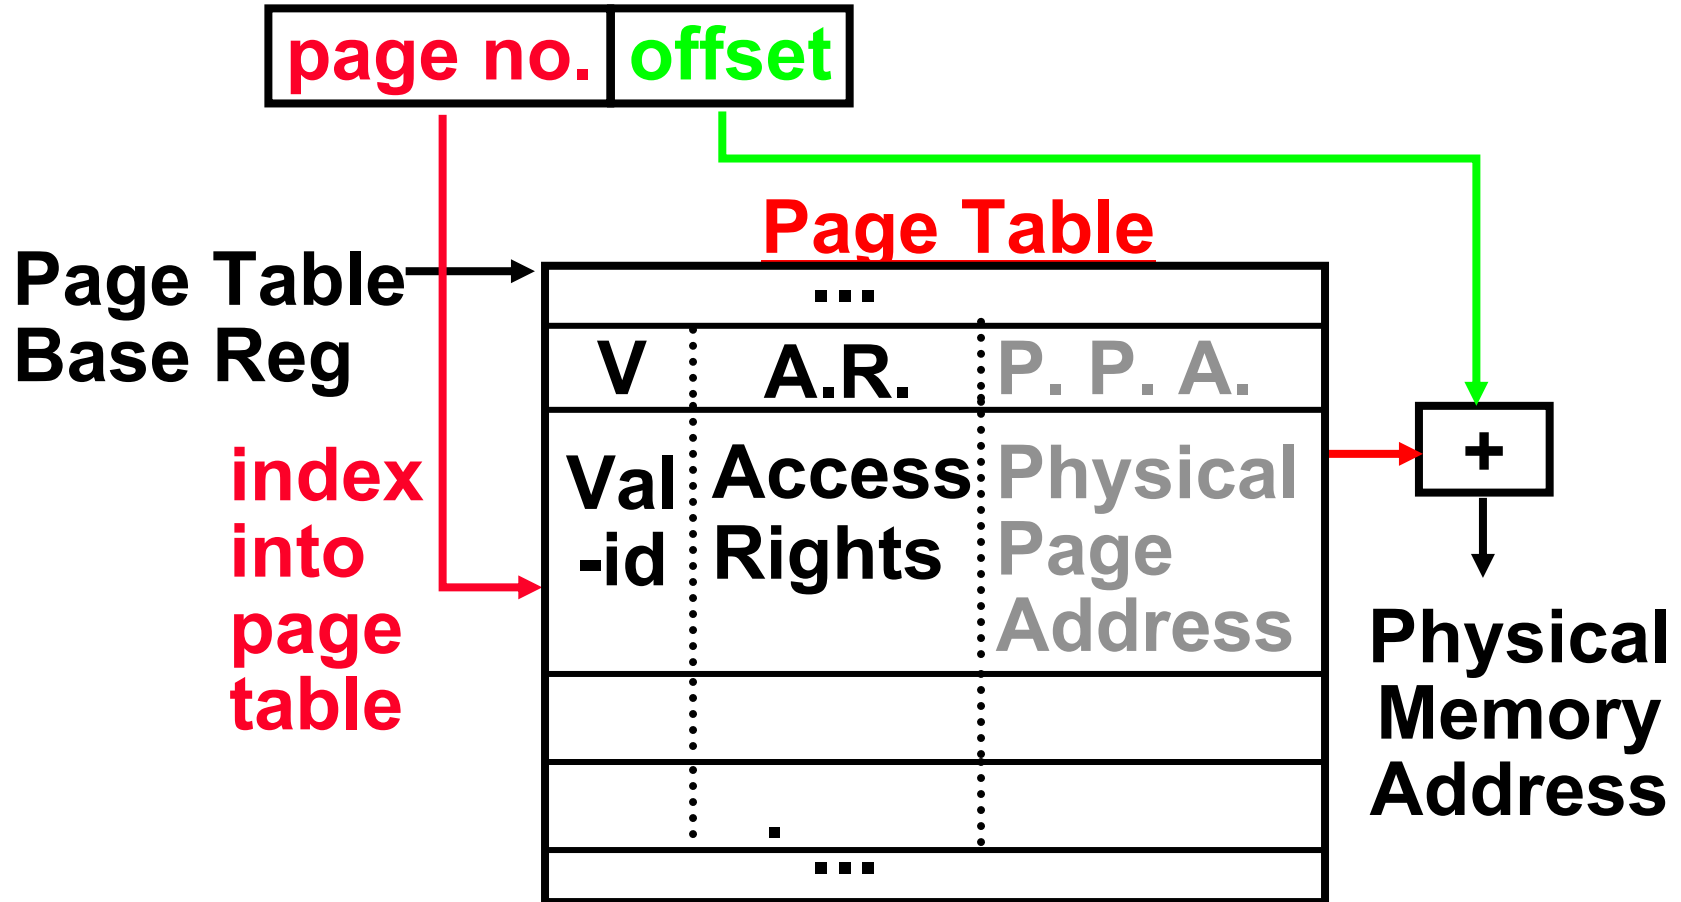
- Each process has its own page table.
- OS changes page tables by changing contents of Page Table Base Register
- Different processes protected from each other
- Kernel data protected from User programs

Sharing:

- Two processes can have a page table entry that is mapped to the same physical page.



Virtual Address:



Page Table located in physical memory


Virtual Memory Problem #1

- Map every address \Rightarrow 1 indirection via Page Table in memory per virtual address \Rightarrow 1 virtual memory accesses = 2 physical memory accesses \Rightarrow SLOW!
- Observation: since locality in pages of data, there must be locality in virtual address translations of those pages
- Since small is fast, why not use a small cache of virtual to physical address translations to make translation fast?
- For historical reasons, cache is called a Translation Lookaside Buffer, or TLB



Translation Look-Aside Buffers

- TLBs usually small, typically 128 - 256 entries
- Like any other cache, the TLB can be direct mapped, set associative, or fully associative



On TLB miss, get page table entry from main memory

What if not in TLB?

- **Option 1: Hardware checks page table and loads new Page Table Entry into TLB**
- **Option 2: Hardware traps to OS, up to OS to decide what to do**
- **MIPS follows Option 2: Hardware knows nothing about page table**



What if the data is on disk?

- **We load the page off the disk into a free block of memory, using a DMA transfer**
 - **Meantime we switch to some other process waiting to be run**
- **When the DMA is complete, we get an interrupt and update the process's page table**
 - **So when we switch back to the task, the desired data will be in memory**



What if we don't have enough memory?

- We chose some other page belonging to a program and transfer it onto the disk if it is dirty
 - If clean (disk copy is up-to-date), just overwrite that data in memory
 - We chose the page to evict based on replacement policy (e.g., LRU)
- And update that program's page table to reflect the fact that its memory moved somewhere else
- If continuously swap between disk and memory, called **Thrashing**



Virtual Memory Problem #2

◦ Not enough physical memory!

- Only, say, 64 MB of physical memory
- N processes, each 4 GB (2^{32} B) of virtual memory!
- Could have 1K virtual pages/physical page!

◦ Spatial Locality to the rescue

- Each page is 4 KB, lots of nearby references

◦ No matter how big program is, at any time only accessing a few pages

- “Working Set”: recently used pages



Virtual Memory Problem #3

◦ Page Table too big!

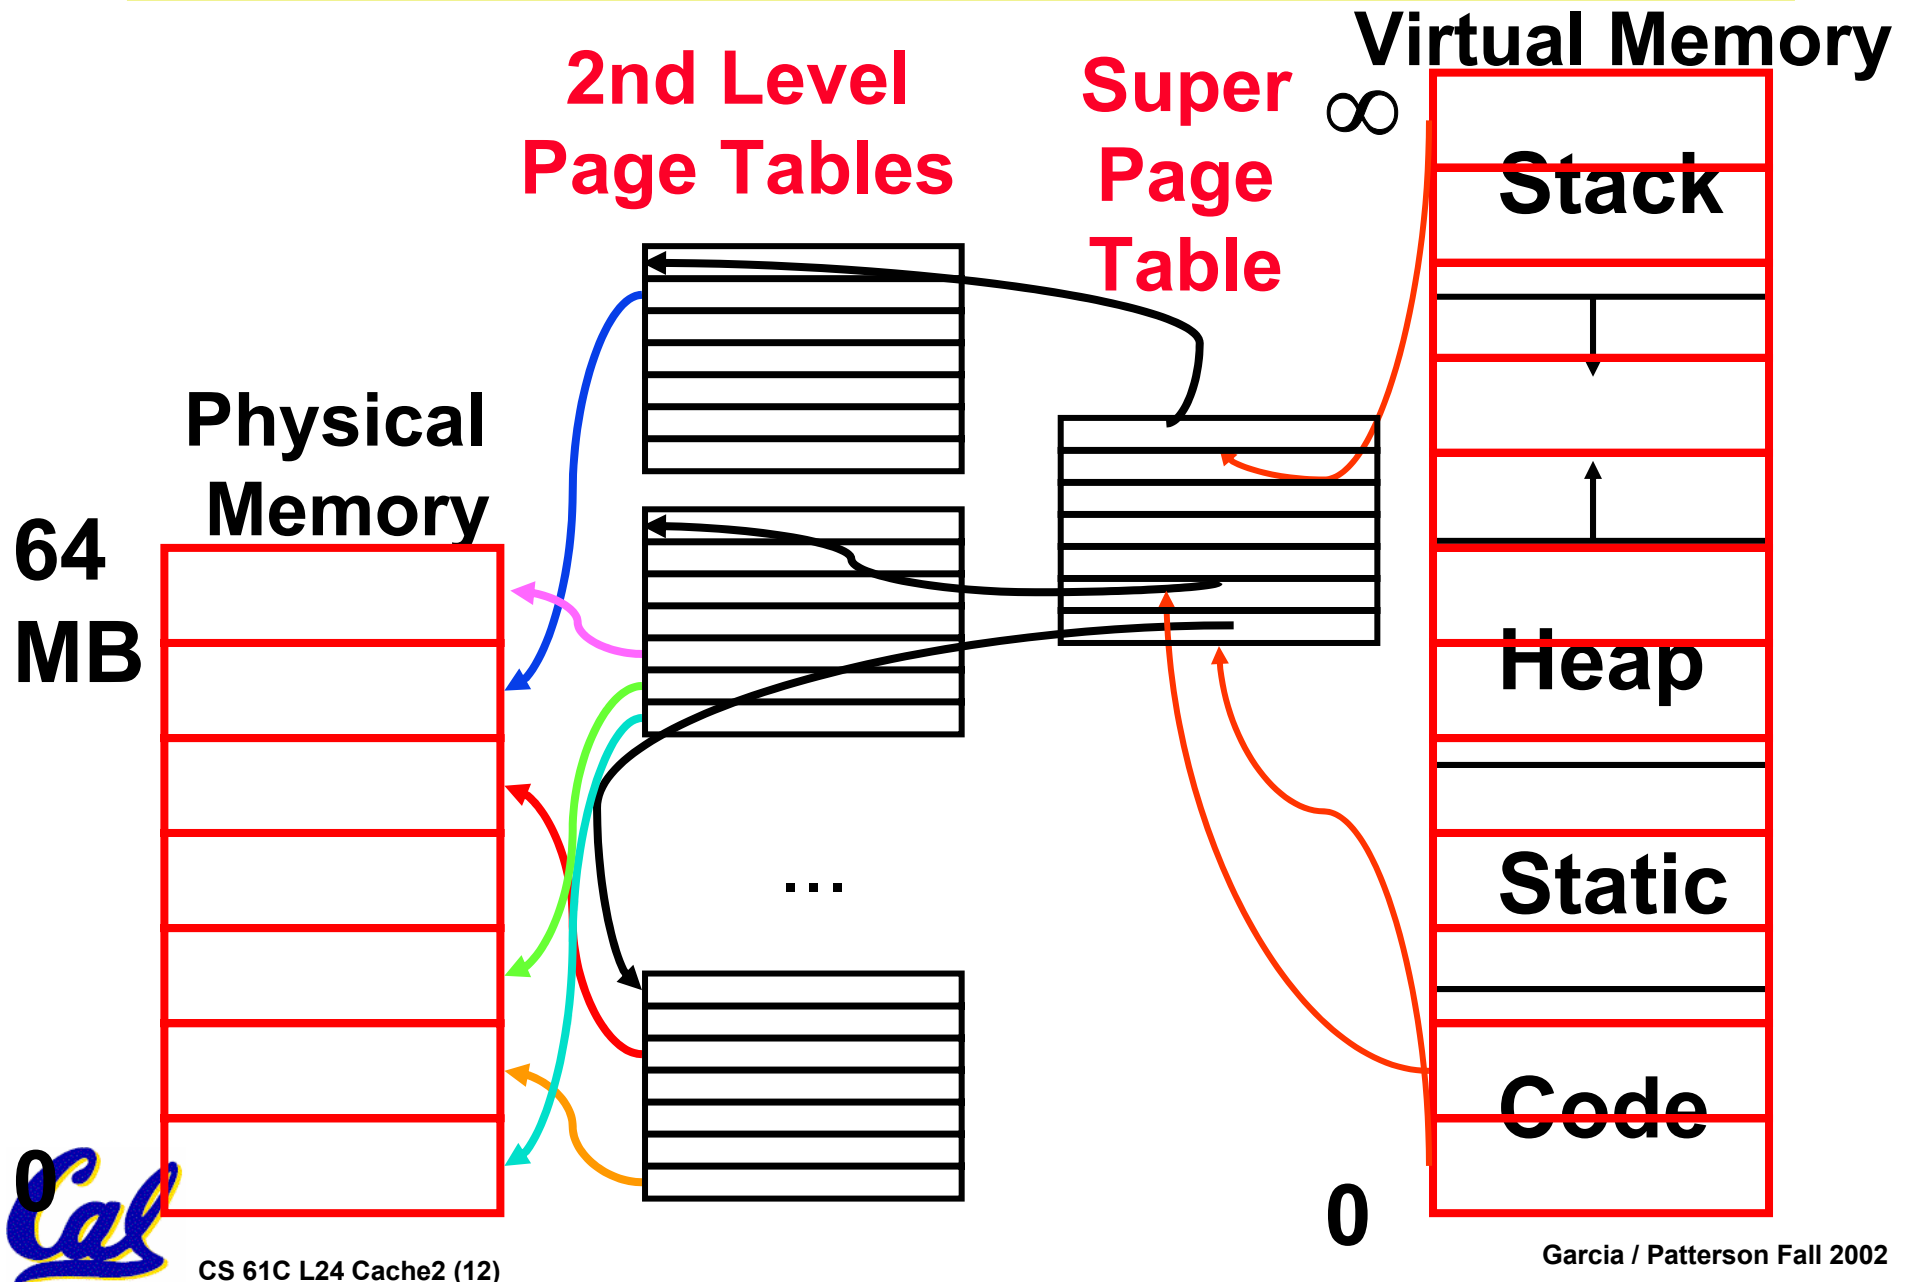
- 4GB Virtual Memory \div 4 KB page
 \Rightarrow ~ 1 million Page Table Entries
 \Rightarrow 4 MB just for Page Table for 1 process,
 25 processes \Rightarrow 100 MB for Page Tables!

◦ Variety of solutions to tradeoff memory size of mapping function for slower when miss TLB

- Make TLB large enough, highly associative so rarely miss on address translation
- CS 162 will go over more options and in greater depth



2-level Page Table



Page Table Shrink :

- **Single Page Table**

Page Number	Offset
--------------------	---------------

20 bits

12 bits

- **Multilevel Page Table**

Super Page No.	Page Number	Offset
---------------------------	------------------------	---------------

10 bits

10 bits

12 bits

- **Only have second level page table for valid entries of super level page table**



Definitions

$$\bullet \text{ performance}(x) = \frac{1}{\text{execution_time}(x)}$$

"X is n times faster than Y" means

$$n = \frac{\text{Performance}(X)}{\text{Performance}(Y)}$$

° CPU execution time for program

= Clock Cycles for a program
x Clock Cycle Time

° = Clock Cycles for a program
Clock Rate



Measuring Time using Clock Cycles (2/2)

° One way to define clock cycles:

Clock Cycles for program

= Instructions for a program
(called “Instruction Count”)

x Average Clock cycles Per Instruction
(abbreviated “CPI”)

° CPI one way to compare two machines
with **same** instruction set, since
Instruction Count would be the same

